

Linear Regression Model

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
In [2]: car_dataset = pd.read_csv("Car_Price.csv")
```

```
In [3]: car_dataset.head()
```

```
Out[3]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	er
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	

5 rows × 26 columns

```
In [4]: car_dataset.shape
```

```
Out[4]: (205, 26)
```

```
In [5]: car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   car_ID                205 non-null    int64  
 1   symboling              205 non-null    int64  
 2   CarName               205 non-null    object  
 3   fueltype              205 non-null    object  
 4   aspiration            205 non-null    object  
 5   doornumber            205 non-null    object  
 6   carbody               205 non-null    object  
 7   drivewheel            205 non-null    object  
 8   enginelocation        205 non-null    object  
 9   wheelbase             205 non-null    float64 
10  carlength             205 non-null    float64 
11  carwidth              205 non-null    float64 
12  carheight             205 non-null    float64 
13  curbweight            205 non-null    int64  
14  enginetype            205 non-null    object  
15  cylindernumber        205 non-null    object  
16  enginesize            205 non-null    int64  
17  fuelsystem            205 non-null    object  
18  boreratio             205 non-null    float64 
19  stroke               205 non-null    float64 
20  compressionratio      205 non-null    float64 
21  horsepower            205 non-null    int64  
22  peakrpm               205 non-null    int64  
23  citympg               205 non-null    int64  
24  highwaympg            205 non-null    int64  
25  price                 205 non-null    float64 
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [6]: car_dataset.describe()
```

Out[6]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	en
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	20
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	12
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	4
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	6
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	9
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	12
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	14
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	32

Extracting Some Features from the Dataset

```
In [7]: car_dataset.CarName.unique()
```

```
Out[7]: array(['alfa-romero giulia', 'alfa-romero stelvio',  
              'alfa-romero Quadrifoglio', 'audi 100 ls', 'audi 100ls',  
              'audi fox', 'audi 5000', 'audi 4000', 'audi 5000s (diesel)',  
              'bmw 320i', 'bmw x1', 'bmw x3', 'bmw z4', 'bmw x4', 'bmw x5',  
              'chevrolet impala', 'chevrolet monte carlo', 'chevrolet vega 2300',  
              'dodge rampage', 'dodge challenger se', 'dodge d200',  
              'dodge monaco (sw)', 'dodge colt hardtop', 'dodge colt (sw)',  
              'dodge coronet custom', 'dodge dart custom',  
              'dodge coronet custom (sw)', 'honda civic', 'honda civic cvcc',  
              'honda accord cvcc', 'honda accord lx', 'honda civic 1500 gl',  
              'honda accord', 'honda civic 1300', 'honda prelude',  
              'honda civic (auto)', 'isuzu MU-X', 'isuzu D-Max ',  
              'isuzu D-Max V-Cross', 'jaguar xj', 'jaguar xf', 'jaguar xk',  
              'maxda rx3', 'maxda glc deluxe', 'mazda rx2 coupe', 'mazda rx-4',  
              'mazda glc deluxe', 'mazda 626', 'mazda glc', 'mazda rx-7 gs',  
              'mazda glc 4', 'mazda glc custom l', 'mazda glc custom',  
              'buick electra 225 custom', 'buick century luxury (sw)',  
              'buick century', 'buick skyhawk', 'buick opel isuzu deluxe',  
              'buick skylark', 'buick century special',  
              'buick regal sport coupe (turbo)', 'mercury cougar',  
              'mitsubishi mirage', 'mitsubishi lancer', 'mitsubishi outlander',  
              'mitsubishi g4', 'mitsubishi mirage g4', 'mitsubishi montero',  
              'mitsubishi pajero', 'Nissan versa', 'nissan gt-r', 'nissan rogue',  
              'nissan latio', 'nissan titan', 'nissan leaf', 'nissan juke',  
              'nissan note', 'nissan clipper', 'nissan nv200', 'nissan dayz',  
              'nissan fuga', 'nissan otti', 'nissan teana', 'nissan kicks',  
              'peugeot 504', 'peugeot 304', 'peugeot 504 (sw)', 'peugeot 604sl',  
              'peugeot 505s turbo diesel', 'plymouth fury iii',  
              'plymouth cricket', 'plymouth satellite custom (sw)',  
              'plymouth fury gran sedan', 'plymouth valiant', 'plymouth duster',  
              'porsche macan', 'porsche panamera', 'porsche cayenne',  
              'porsche boxster', 'renault 12tl', 'renault 5 gtl', 'saab 99e',  
              'saab 99le', 'saab 99gle', 'subaru', 'subaru dl', 'subaru brz',  
              'subaru baja', 'subaru rl', 'subaru r2', 'subaru trezia',  
              'subaru tribeca', 'toyota corona mark ii', 'toyota corona',  
              'toyota corolla 1200', 'toyota corona hardtop',  
              'toyota corolla 1600 (sw)', 'toyota carina', 'toyota mark ii',  
              'toyota corolla', 'toyota corolla liftback',  
              'toyota celica gt liftback', 'toyota corolla tercel',  
              'toyota corona liftback', 'toyota starlet', 'toyota tercel',  
              'toyota cressida', 'toyota celica gt', 'toyota tercel',  
              'volkswagen rabbit', 'volkswagen 1131 deluxe sedan',  
              'volkswagen model 111', 'volkswagen type 3', 'volkswagen 411 (sw)',  
              'volkswagen super beetle', 'volkswagen dasher', 'vw dasher',  
              'vw rabbit', 'volkswagen rabbit', 'volkswagen rabbit custom',  
              'volvo 145e (sw)', 'volvo 144ea', 'volvo 244dl', 'volvo 245',  
              'volvo 264gl', 'volvo diesel', 'volvo 246'], dtype=object)
```

```
In [8]: car_dataset.price.sum()
```

```
Out[8]: 2721725.667
```

```
In [9]: car_dataset.carbody.unique()
```

```
Out[9]: array(['convertible', 'hatchback', 'sedan', 'wagon', 'hardtop'],  
              dtype=object)
```

```
In [10]: car_dataset.isnull().sum()
```

```
Out[10]: car_ID          0
symboling              0
CarName                0
fueltype               0
aspiration             0
doornumber             0
carbody                0
drivewheel             0
enginelocation         0
wheelbase              0
carlength              0
carwidth               0
carheight              0
curbweight             0
enginetype             0
cylindernumber         0
enginesize             0
fuelsystem             0
boreratio              0
stroke                 0
compressionratio       0
horsepower             0
peakrpm                0
citympg                0
highwaympg             0
price                  0
dtype: int64
```

```
In [11]: print(car_dataset.fueltype.value_counts())
print(car_dataset.carbody.value_counts())
```

```
gas          185
diesel       20
Name: fueltype, dtype: int64
sedan        96
hatchback    70
wagon        25
hardtop       8
convertible   6
Name: carbody, dtype: int64
```

Training the Model

```
In [12]: y=car_dataset[["price"]]
x=car_dataset[["enginesize"]]
```

```
In [13]: lm = LinearRegression()
```

```
In [15]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=101)
```

```
In [16]: lm.fit(x_train,y_train)
```

```
Out[16]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [17]: print(lm.intercept_,lm.coef_)  
[-7107.69540612] [[161.34743434]]
```

Predicting the Value After the Training Phase

```
In [18]: pred=lm.predict(x_test)  
print(pred[0:10])
```

```
[[ 7413.57368461]  
 [10479.1749371 ]  
 [20482.71586628]  
 [ 8704.35315935]  
 [12253.99671486]  
 [22096.19020969]  
 [12415.3441492 ]  
 [ 8543.005725  ]  
 [10640.52237144]  
 [ 7736.2685533 ]]
```


```
In [19]: training_pred=lm.predict(x_test)
```

```
In [20]: error_score=metrics.r2_score(y_test,training_pred)  
print(error_score)
```

```
0.8258154601020361
```


Plotting Some Graphs after the actual and the predicted Values

```
In [21]: plt.scatter(y_test,training_pred,color="red")  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.show()
```



```
In [22]: sns.jointplot(y=car_dataset[["price"]],x=car_dataset[["enginesize",]],data=car_
dataset,kind="reg",color="blue")
```


```
Out[22]: <seaborn.axisgrid.JointGrid at 0x11b9f445488>
```



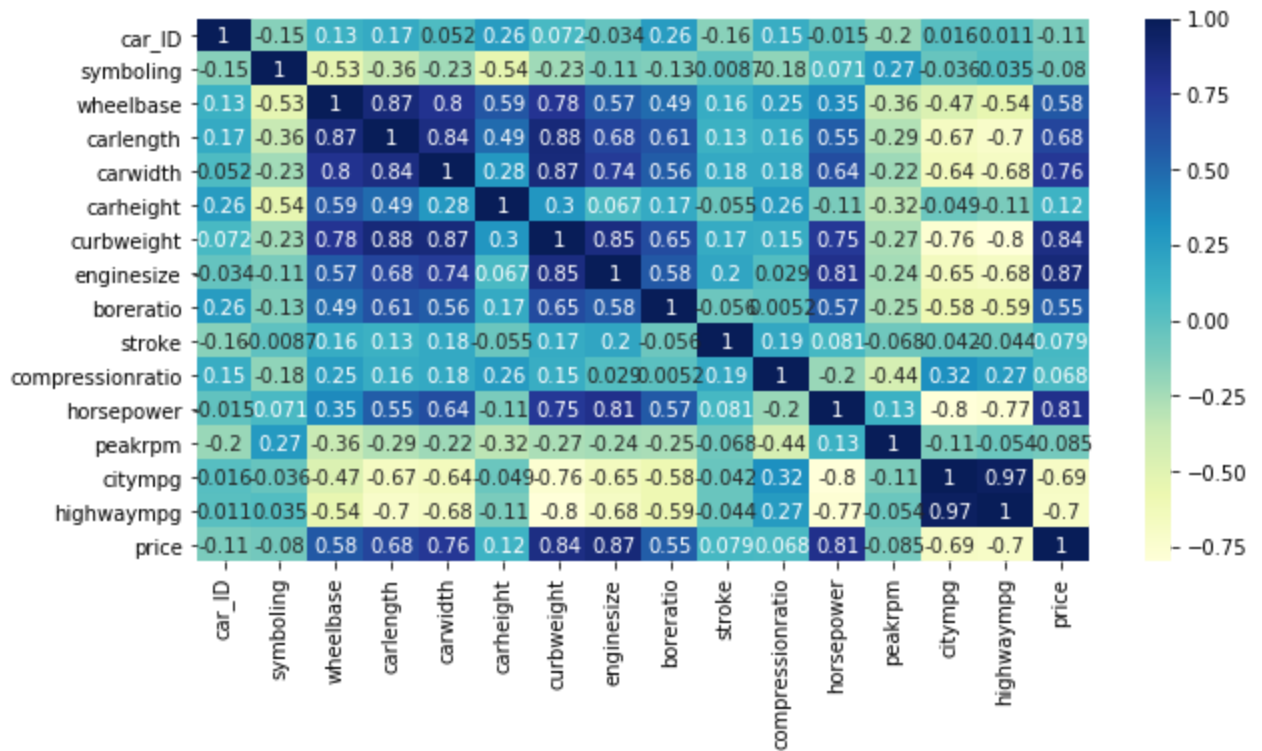
Plotting Some Graphs related to the Dataset

```
In [23]: sns.distplot(y_test-pred,bins=50,color="red",kde=False)
```


```
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x11b9f665d48>
```




```
In [24]: plt.figure(figsize = (10, 5))
sns.heatmap(car_dataset.corr(), annot = True, cmap="YlGnBu")
plt.show()
```



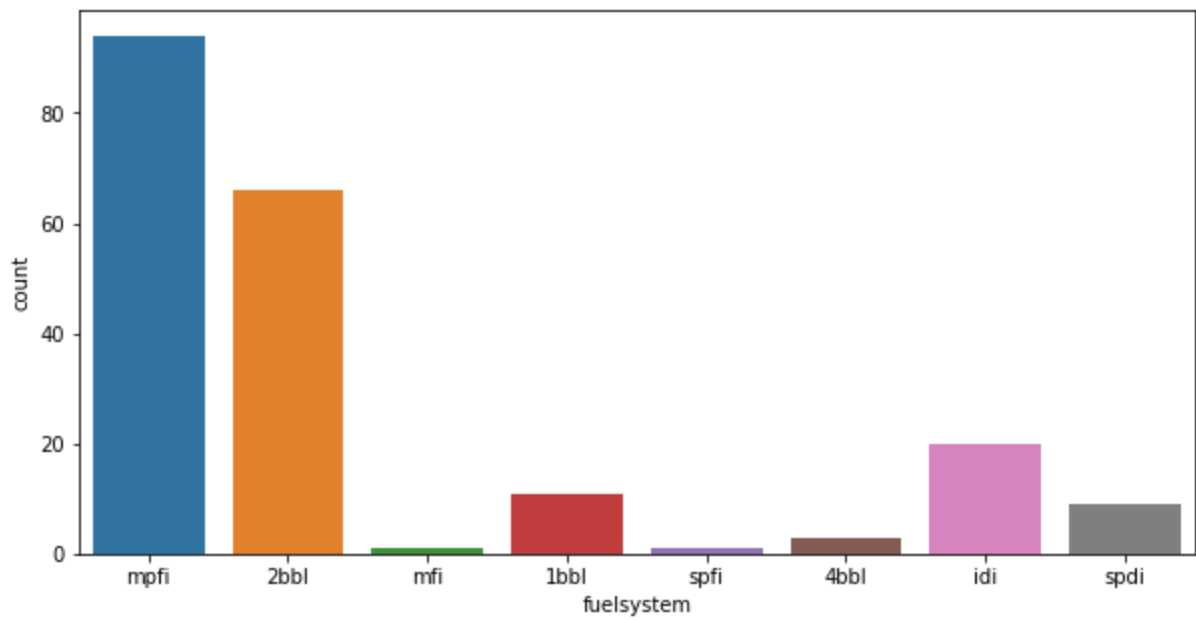
```
In [25]: plt.figure(figsize = (10,5))
sns.countplot(x="enginesize",data=car_dataset)
plt.show()
```



```
In [26]: plt.figure(figsize = (10,5))  
sns.violinplot(x="drivewheel",y="price",data=car_dataset)  
plt.show()
```

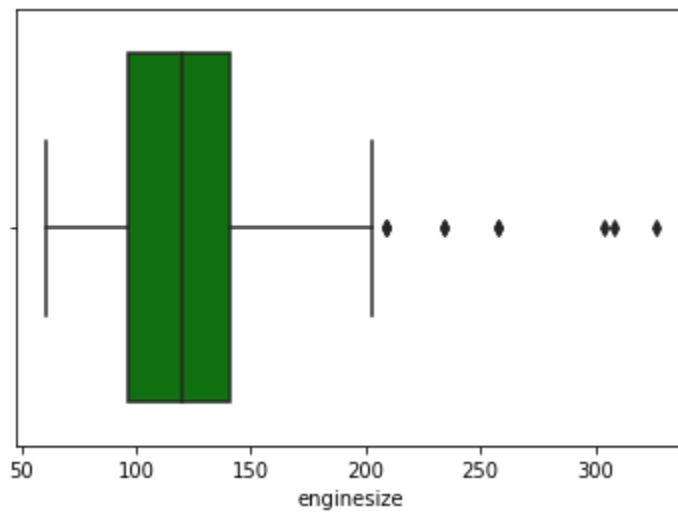


```
In [27]: plt.figure(figsize = (10,5))  
sns.countplot(x="fuelsystem",data=car_dataset)  
plt.show()
```





```
In [28]: sns.boxplot(x="engine_size",data=car_dataset,color="green")
```

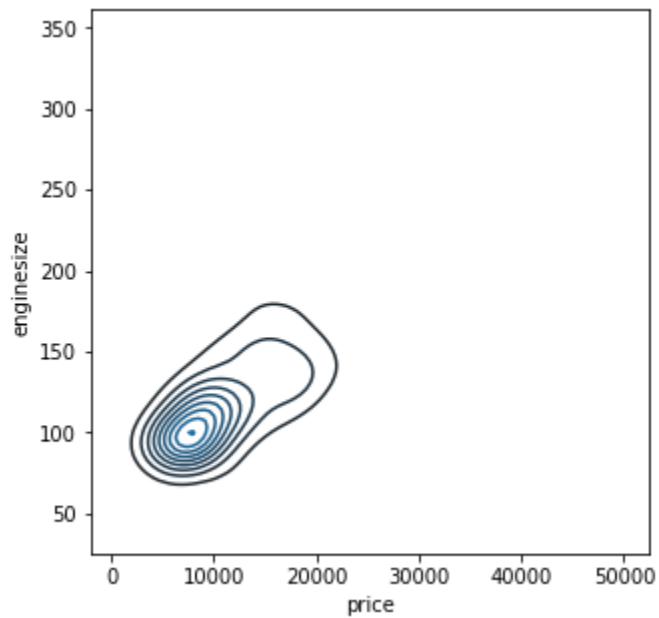
```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x11b9fc27248>
```



```
In [29]: plt.figure(figsize = (5,5))  
sns.stripplot(x="engine_size",y="price",data=car_dataset)  
plt.show()
```



```
In [30]: plt.figure(figsize = (5,5))
sns.kdeplot(car_dataset.price,car_dataset.enginesize)
plt.show()
```



```
In [31]: df= pd.DataFrame(car_dataset.groupby(['fueltype'])['price'].mean().sort_values(
ascending = False))
df.plot.bar()
plt.title('Fuel Type vs Average Price')
plt.show()
```

